High-Level Risk Assessment of X's Community Notes:

The Boundary Conditions of Crowdsourced Correction

Dongjae Kang

TPINIA7008
Principles & Practice of Online Trust & Safety

Professor Tim Bernard

Spring 2026

Columbia University
School of International and Public Affairs

High-Level Risk Assessment of X's Community Notes:

The Boundary Conditions of Crowdsourced Correction

by Dongjae Kang

X's Community Notes, launched as Birdwatch in 2021 and rebranded in 2022, allows approved contributors to attach context or fact-check style notes to posts on the platform. Since X wound down centralized fact-checking, Community Notes has become the platform's primary correction mechanism. Recent empirical research suggests the system reduces engagement with false content: attaching a note has been associated with roughly 46% fewer reposts and 44% fewer likes within 48 hours, and posts that receive a public note are about 32% more likely to be deleted by their authors.

The approach is moving beyond X. On January 7, 2025, Meta announced that it would end its third-party fact-checking program in the United States and adopt a Community Notes model on Facebook, Instagram, and Threads. YouTube has been piloting a Notes feature since June 2024, initially on mobile English content in the United States. Crowdsourced correction, in other words, is becoming a default design pattern for misinformation governance across major platforms, which makes careful attention to the feature's boundary conditions more, not less, important.

A risk assessment is not a claim that an intervention has failed. It asks under what conditions the intervention performs as intended, and where gaps remain that design attention can address. This assessment maps risks associated with Community Notes through two stylized user types positioned at opposite ends of engagement with political content, applies the eSafety Commissioner's Safety by Design principles — Service Provider Responsibility, User Empowerment and Autonomy, and Transparency and Accountability — and proposes recommendations intended to complement the feature rather than replace it. One lens used in the gap analysis is a *provenance view*, which asks how the feature treats the question of where information originates and how its authority is established before correction is contemplated. Each identified gap is accompanied by severity and likelihood notations, and each recommendation is paired with a trade-off against user experience and a practical implementation challenge.

User Journey and Identified Risks

User Type and Journey	Identified Risks	Severity	Potential Types of Harm
The Low-Engagement Encounter. A user who opens X during leisure hours, scrolls through political posts in their feed, and rarely taps into embedded notes or expanded context. In the 24-48 hour window during which a misleading post typically accumulates most of its virality, this user forms an initial	The correction signal may never register: notes attached more than ~48 hours after publication show limited effect on engagement and belief. Low salience of the note UI on mobile feeds means exposure does not imply reading. Virality diffusion on polarizing topics often outpaces bridging-	High	**Personal:** individual belief formation and civic judgment. **Community / Societal:** aggregated civic impact across large numbers of voters during election cycles.

impression well before a bridging-consensus note is attached, if one is attached at all.	consensus formation, leaving the most consequential window uncorrected.		
The High-Engagement Partisan Rater. A politically active user who spends substantial time on X, reads posts across perspectives, and contributes ratings that partly determine which notes earn public display. Motivated reasoning and group identity shape rating behavior on contested topics.	Partisanship has been shown to be a stronger predictor of rating behavior than content accuracy on polarizing domains. Roughly 30% of displayed notes later lose their "helpful" status. A small coalition of 5-20% of coordinated bad-faith raters can strategically suppress targeted helpful notes.	Medium–High	**Personal:** cognitive load and motivated-reasoning strain on the rater. **Structural / Ecosystem:** rating inputs shape which corrections the broader public ever sees.

Existing Safety Gaps

Attachment Latency. *(Severity: High | Likelihood: High for viral political content.)* Community Notes' effectiveness depends heavily on how quickly a note attaches. Recent PNAS research finds that the reduction in reposts, likes, and views is concentrated in the early hours after attachment, and that notes arriving more than roughly 48 hours after publication show limited measurable impact. For the most viral misleading posts, the window in which correction can meaningfully reshape the information environment is short.

Coverage Ceiling. *(Severity: High | Likelihood: High on polarized topics.)* Among notes submitted by contributors, fewer than 10% are ever publicly displayed. The bridging algorithm's requirement — favorable ratings from contributors whose past rating patterns have differed — is hardest to satisfy on the most polarizing topics. The correction mechanism is least available precisely where correction is most consequential. The same consensus threshold also creates an opening for misuse: a coordinated coalition of roughly 5-20% of bad-faith raters has been shown to be able to strategically suppress helpful notes on targeted topics, amplifying the coverage gap on exactly the content categories where risk is highest.

Consensus Instability. *(Severity: Medium | Likelihood: Medium.)* Approximately 30.2% of displayed notes later lose "helpful" status and are removed. A user exposed to the post while a note was attached sees the correction; a later viewer may not. The state of the information environment at any moment is a function not only of what is posted, but of which notes are currently inside or outside consensus.

Salience and Cognitive Asymmetry. *(Severity: Medium | Likelihood: High at both extremes of engagement.)* Research on note-attachment effects finds that attaching a note to a low-engagement post can paradoxically increase the post's visibility, while at the high-engagement end, users process many signals simultaneously and may not attend to note content in proportion to its display. The feature's effectiveness is contingent on reading behavior that is unevenly distributed across the user base and poorly instrumented.

Provenance in the Correction Signal. *(Severity: Medium | Likelihood: Persistent / structural.)* Community Notes operates on the text of contributor-written notes rather than on verifiable signals about the origin of the underlying content. A desk-verified legacy media report and an uncited assertion both enter the bridging algorithm as note text, with authority absorbed only indirectly through what raters find persuasive. This is structurally distinct from provenance infrastructures — such as cryptographically signed content credentials — that mark origin before distribution. Whether treating source authority as an unmodeled input is a gap or a deliberate choice is open to discussion; the point for a risk assessment is that the current architecture leaves origin verification outside the correction signal.

Cross-Context Portability. *(Severity: Medium–High | Likelihood: Increases as Meta and YouTube extensions scale.)* As Meta extends Community Notes to Facebook, Instagram, and Threads, and as YouTube pilots its version on video, the assumption that bridging will function comparably in each context is not automatic. Conditions for productive collective authoring vary across linguistic, platform, and cultural settings. The comparatively slow development of Korean Wikipedia — which, despite Korea's substantial internet infrastructure, has grown more slowly than peers and operates alongside competing crowd-authored and expert-curated alternatives — illustrates that crowd-authored information systems face uneven institutional and cultural conditions outside English-language ecosystems. This is not a prediction of failure but a reminder that success is conditional on local dynamics that warrant pre-deployment attention.

Key Recommendations

1. Accelerate note attachment without foreclosing deliberation.

- Evaluate hybrid pipelines that pair platform-side pre-flagging of high-virality claims with human-in-the-loop bridging validation, shortening the time from publication to attachment.
- Prioritize latency reduction on posts whose predicted virality exceeds a threshold, given evidence that notes arriving after ~48 hours show limited effect.
- Avoid fully automated note generation, which would reintroduce the centralized model Community Notes was designed to move beyond.
- Publish the latency-to-attachment distribution, disaggregated by topic category, to enable external assessment.

This recommendation engages the Safety by Design principle of ***Service Provider Responsibility****: the burden of catching fast-moving misinformation should not rest entirely on a volunteer crowd whose consensus-formation is, by design, slow. A practical measurement is median time-to-display by topic bucket.*

Trade-off. Platform-side pre-flagging partially reintroduces the centralized intervention dynamic Community Notes was designed to decentralize. Transparency about what triggers pre-flagging — and audit access to its inputs — is required to preserve the feature's legitimacy as a crowdsourced mechanism.

Implementation challenge. Calibrating virality-prediction thresholds without inducing over-flagging that overwhelms rater capacity is non-trivial. Misspecification risks either false-positive corrections on legitimate content or continued latency on the most viral misleading posts.

2. Design for the full engagement distribution.

- Treat the low-engagement user as a primary design constraint: note salience, pre-click summary text, and visual trust cues should not assume the user will tap to read.
- For high-engagement raters, consider cognitive-load instrumentation (e.g., session-level rating volume feedback) that preserves judgment quality as participation scales.
- Report publicly what share of users exposed to a flagged post read note content, not only the share of flagged posts receiving notes.
- Investigate whether note attachment on low-initial-engagement posts has produced unintended visibility effects, and adjust attachment rules accordingly.

This recommendation engages the Safety by Design principle of ***User Empowerment and Autonomy****: correction design should support comprehension at the point of exposure, without demanding additional effort from users who came to the platform for other reasons. Effectiveness can be measured through exposure-to-read ratios rather than display counts alone.*

Trade-off. Richer note UI elements add friction to a scrolling feed, which conflicts with platform incentives that optimize for session time and engagement velocity. More assertive correction signals may also be perceived as paternalistic by users who prefer minimal interface intrusion, creating a tension between safety salience and user autonomy.

Implementation challenge. Instrumenting exposure-to-read metrics reliably requires on-device attention telemetry whose collection raises privacy concerns, and whose validation is difficult without parallel access to user behavior data that platforms are increasingly reluctant to share.

3. Open rater composition and consensus criteria to independent audit, and treat cross-platform extension as a research question.

- Publish aggregate rater composition (activity-volume distributions, jurisdictional spread, language capacity) in forms accessible to external researchers, while protecting individual privacy.
- Create data-sharing arrangements with civil-society and academic auditors for study of bridging outcomes on polarized and cross-lingual topics.
- For Meta's and YouTube's extensions, invest in pre-deployment research on local collective-authoring conditions — trust-graph density, coordinated-rater risk, institutional substitutes — rather than porting the English-language model by default.
- Consider complementary provenance infrastructures — such as signed-origin metadata applied at the point of content creation — as a structurally upstream layer that operates independently of, and could usefully support, crowd-based correction downstream.

This recommendation engages the Safety by Design principle of ***Transparency and Accountability****, and recognizes that the feature's portability across platforms and languages is a governance question, not only a product question. A concrete measure is the number of independent audits published per year with platform cooperation, alongside auditor-reported bridging-outcome disparities across languages.*

Trade-off. Publishing rater-composition data risks enabling coordinated targeting of rater communities by adversaries, placing transparency and anti-manipulation goals in direct tension. Careful aggregation (thresholded bucket sizes, differential-privacy guarantees) is required so that accountability does not come at the cost of rater safety.

Implementation challenge. Cross-context pre-deployment research delays product rollout, which platforms under competitive pressure are unlikely to prioritize absent external regulatory or reputational accountability. Complementary provenance infrastructures, meanwhile, require cross-industry coordination and standards adoption that no single platform can secure alone.

Community Notes has compiled empirical evidence of effectiveness that few content-moderation interventions can match, and its adoption by Meta signals that crowdsourced correction is becoming the default platform-governance approach to misinformation. The boundary conditions identified here — attachment latency, coverage ceiling, consensus instability, salience asymmetry, an unmodeled provenance layer, and uneven cross-context portability — do not undermine that evidence. They suggest where continued design, transparency, and complementary infrastructure could extend the feature's reach without displacing its crowdsourced character.

References

1. Slaughter, I., Peytavin, A., Ugander, J., & Saveski, M. (2025). Community notes reduce engagement with and diffusion of false information online. *Proceedings of the National Academy of Sciences, 122*(38), e2503413122. https://www.pnas.org/doi/10.1073/pnas.2503413122

2. Community Notes. (n.d.). *Wikipedia.* Retrieved April 2026 from https://en.wikipedia.org/wiki/Community_Notes

3. Community Notes are Vulnerable to Rater Bias and Manipulation. (2025). *arXiv preprint arXiv:2511.02615.* https://arxiv.org/html/2511.02615

4. Consensus Stability of Community Notes on X. (2026). *arXiv preprint arXiv:2601.14002.* https://arxiv.org/html/2601.14002

5. Can Crowdchecking Curb Misinformation? Evidence from Community Notes. (2025). *Information Systems Research.* https://pubsonline.informs.org/doi/10.1287/isre.2024.1609

6. Meta. (2025, January 7). More Speech and Fewer Mistakes. *Meta Newsroom.* https://about.fb.com/news/2025/01/meta-more-speech-fewer-mistakes/

7. YouTube. (2024, June). Testing new ways to offer viewers more context and information on videos. *YouTube Official Blog.* https://blog.youtube/news-and-events/new-ways-to-offer-viewers-more-context/

8. eSafety Commissioner. (n.d.). Safety by Design. *Australian Government.* https://www.esafety.gov.au/industry/safety-by-design

9. Coalition for Content Provenance and Authenticity (C2PA). (n.d.). *C2PA.* https://c2pa.org/

10. Korean Wikipedia. (n.d.). *Wikipedia.* Retrieved April 2026 from https://en.wikipedia.org/wiki/Korean_Wikipedia